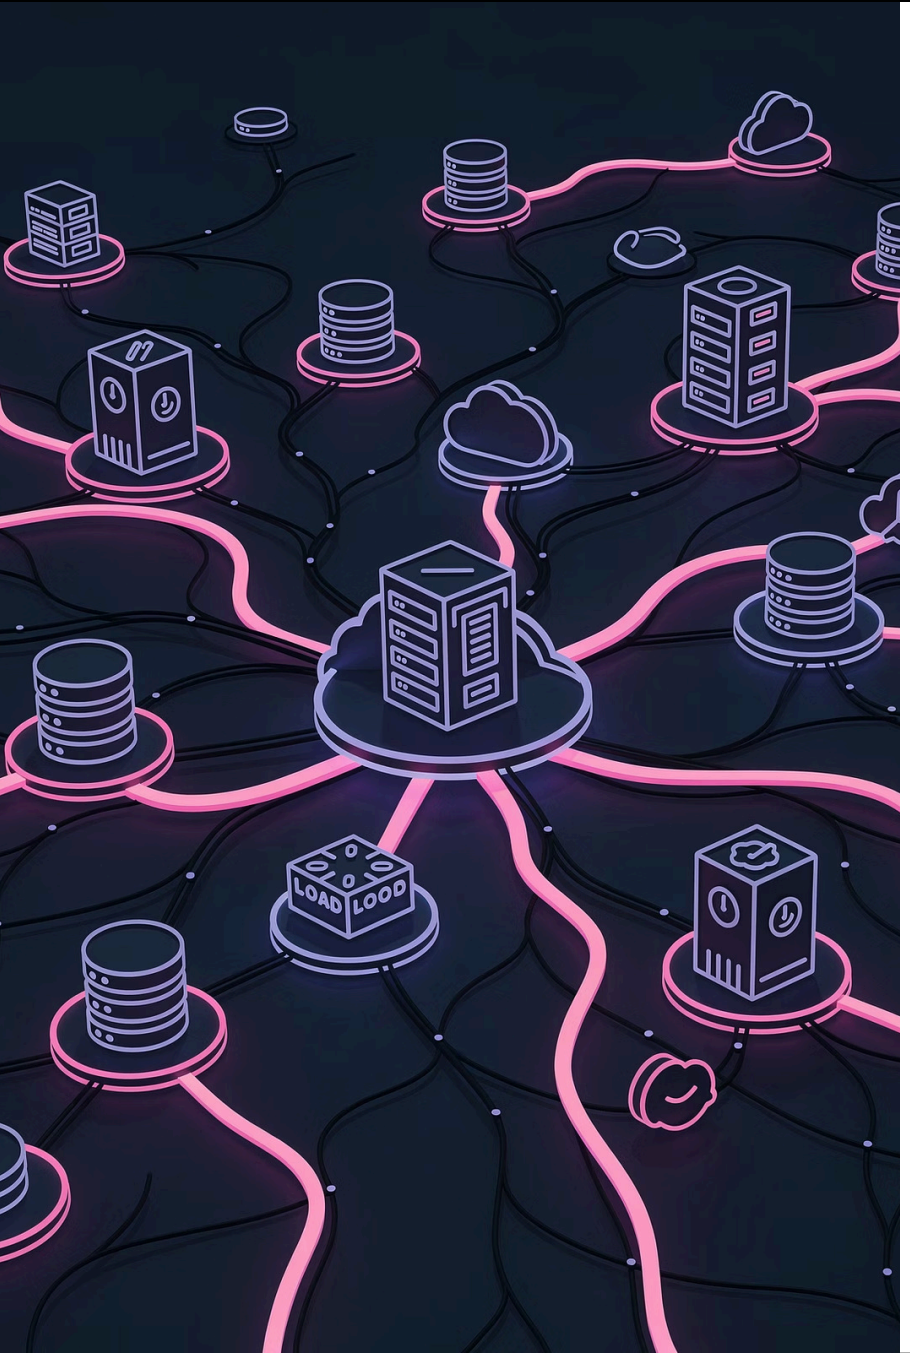

GenAI on AWS

AWS GenAI Service Landscape · Current State, February 2026

A technical deep-dive into Amazon Bedrock, SageMaker JumpStart, foundation models, and the new 2026 service tiers — everything you need to make informed GenAI architecture decisions on AWS.

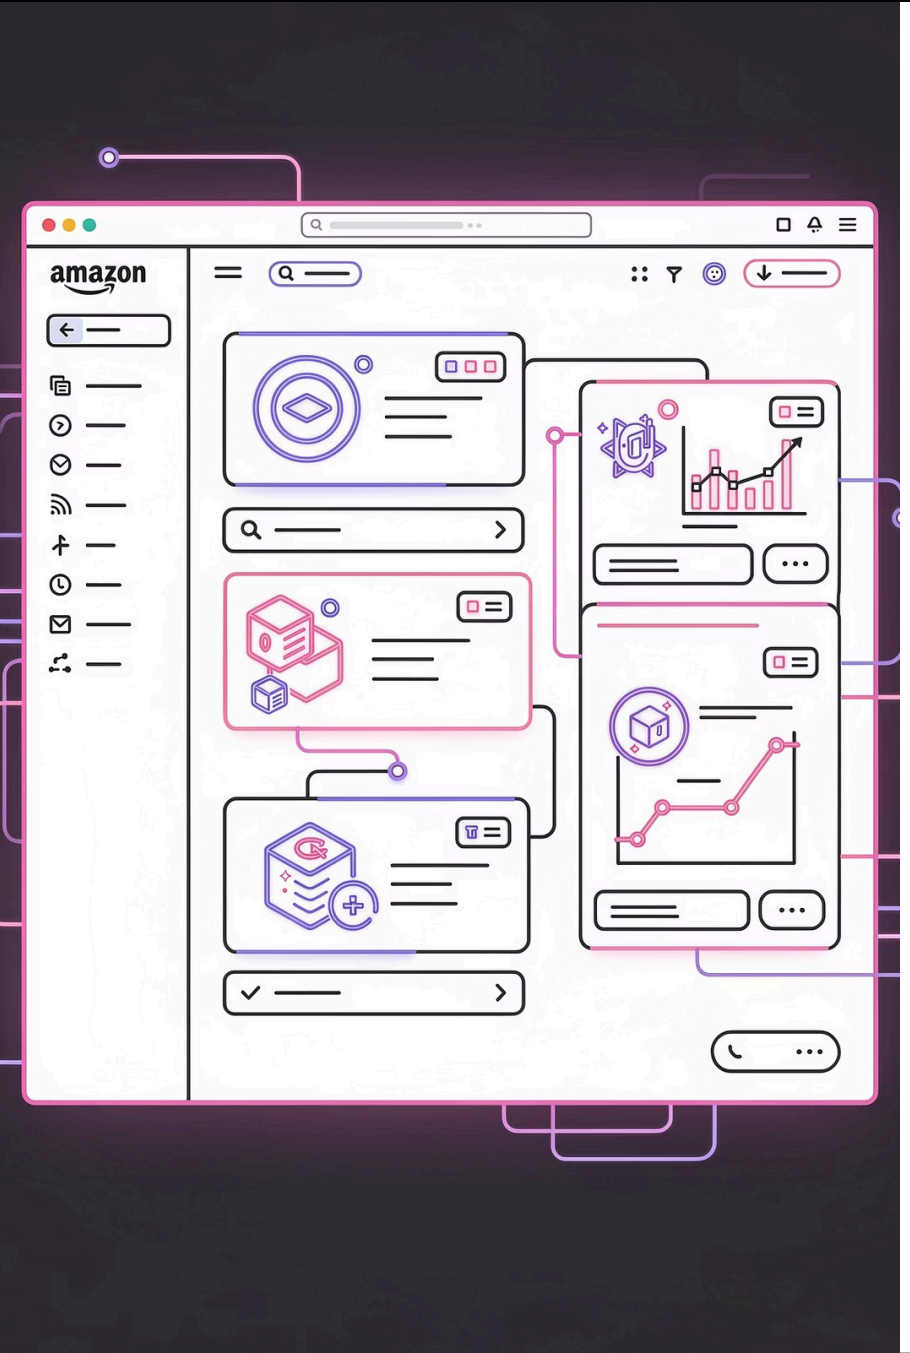

Amazon Bedrock

A **fully managed service** giving access to **100+ foundation models** from leading AI providers — no infrastructure management required. Data is never shared with model providers, with encryption in transit and at rest.

Model Providers

Amazon, Anthropic, OpenAI, Meta, Mistral AI, DeepSeek, AI21 Labs, Cohere, Stability AI, Qwen, MiniMax, Moonshot AI, Writer, and more. **docs.aws.amazon**

2026 Additions

OpenAI gpt-oss-120b & gpt-oss-20b, DeepSeek V3.2, MiniMax M2.1, GLM 4.7, Kimi K2.5, Qwen3 Coder Next. **docs.aws.amazon**

Enterprise Security

Built-in compliance with GDPR, HIPAA, SOC, and FedRAMP High — all via a unified API across every model. **docs.aws.amazon**

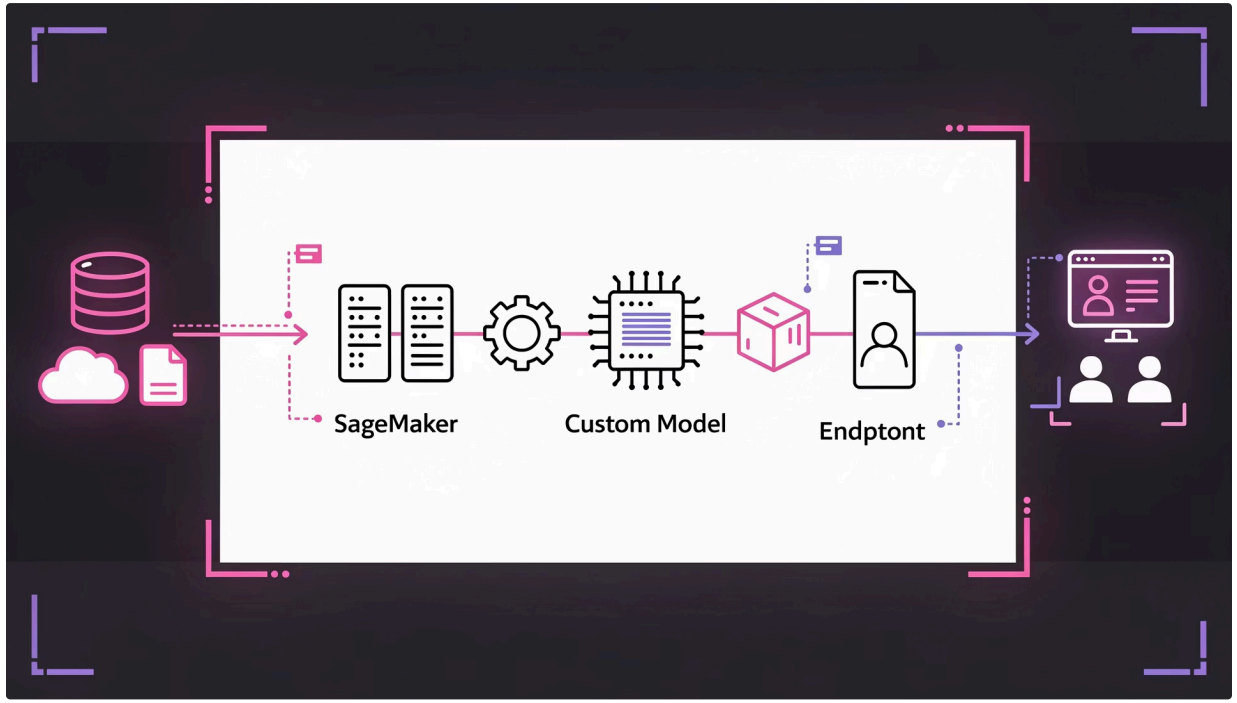

ML DEEP DIVE

Amazon SageMaker JumpStart

Deploy and fine-tune open-source and proprietary models with **full infrastructure control** — ideal for teams with deep ML expertise who need to go beyond managed defaults.

Key Differentiators

- Full control over model deployment and compute
- Custom training and extensive fine-tuning
- Purpose-built ML pipelines beyond GenAI

Bedrock Integration

Models trained in JumpStart can be **registered directly with Bedrock**, unlocking Knowledge Bases, Guardrails, and the Converse API. **docs.aws.amazon**

When to Use Each Service

Choose Amazon Bedrock

- Rapid production deployment without deep ML expertise
- Managed infrastructure with automatic scaling
- Enterprise-grade security and compliance out-of-the-box
- Multi-model experimentation and evaluation
- RAG systems, agents, and conversational applications

Choose SageMaker JumpStart

- Full control over deployment and infrastructure
- Extensive custom model training and fine-tuning
- Specialized ML requirements beyond managed services
- Cost optimization via custom infrastructure
- Complex ML pipelines that extend beyond GenAI

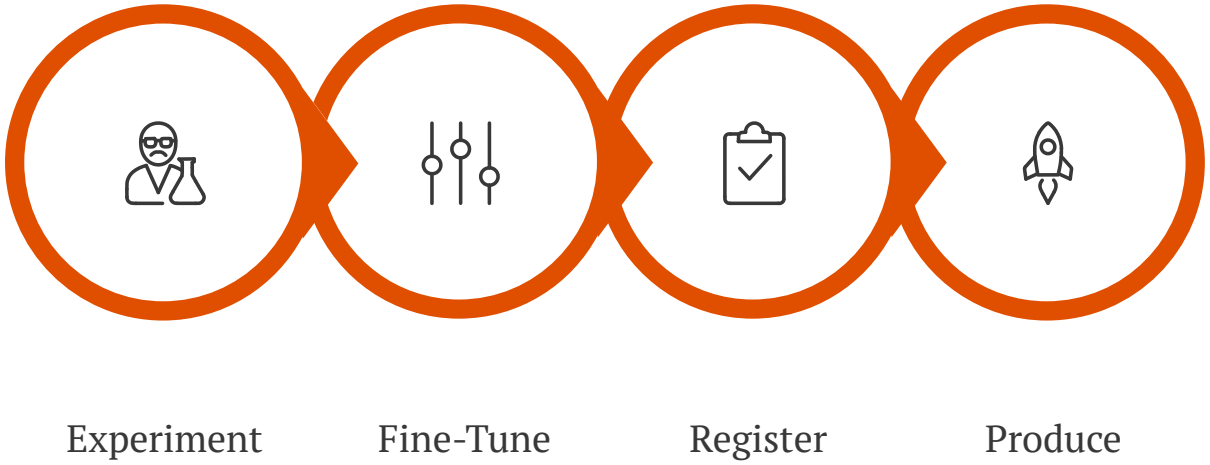

Use the services together: fine-tune custom models in JumpStart, then register them with Bedrock for production-grade managed features and consistent API access.

NEW IN 2026

Amazon Bedrock Marketplace

Discover, test, and subscribe to **100+ additional foundation models** beyond the core Bedrock catalog — all from a unified model discovery interface.

- Seamless integration with Bedrock Agents, Knowledge Bases, and Guardrails
- Managed endpoints with configurable scalability
- Expanded providers including Zendesk, IBM, and more

aws.amazon.com/bedrock/marketplace

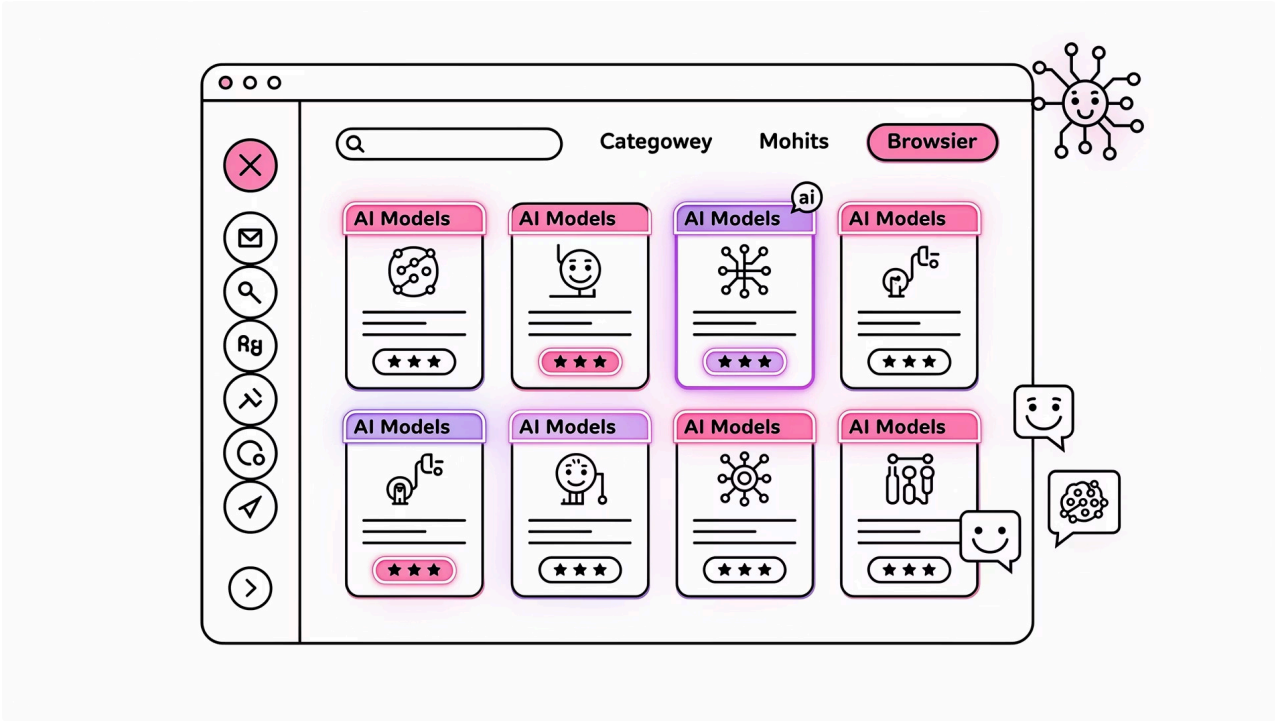

Foundation Models in Bedrock: Flagship Providers

Amazon Nova

Multimodal (text, image, video, documents) with Nova Pro & Premier for extended RAG workflows. **aboutamazon**

Anthropic Claude 4.6

Latest Opus, Sonnet, and Haiku variants. Advanced reasoning, long context windows. Sonnet available in Reserved tier. **aboutamazon**

OpenAI (OSS)

gpt-oss-120b (117B params) and gpt-oss-20b — Mixture-of-Experts architecture, 128K context, Apache 2.0 license. **aws.amazon**

Foundation Models: Extended Provider Ecosystem

Meta Llama Family

Llama 3 through 4 series — 7B to 70B+ parameters, conversational and chat-optimized variants. **aws.amazon**

Mistral AI

Mistral 7B general-purpose, Mistral 7B Instruct for chat, plus newer Ministral and Magistral models. **docs.aws.amazon**

DeepSeek V3.2

Advanced reasoning and agentic intelligence, optimized for complex multi-step workflows. **cloudprice**

Coding Specialists

MiniMax M2.1 for autonomous coding, Qwen3 Coder Next for code generation, Kimi K2.5 for reasoning tasks. **docs.aws.amazon**

AI21, Cohere, Stability

Jamba 1.5 Mini/Large for language tasks; Command & Embed for retrieval; Stable Diffusion for image generation. **docs.aws.amazon**

GLM 4.7 & Flash

Cost-efficient models engineered for production deployment at scale. **docs.aws.amazon**

Bedrock Service Tiers (New in 2026)

Four tiers let you match every workload to the right balance of performance, availability, and cost. aws.amazon

Reserved

Guaranteed capacity, 99.5% uptime SLA, 1–3 month commitment. Fixed token-per-minute pricing for mission-critical apps.

Priority

Up to **25% better OTPS latency** with preferential processing. No reservation required. Ideal for customer-facing chatbots.

Standard

Default tier. Balanced cost and performance with moderate throughput guarantees. Best for everyday content generation tasks.

Flex

Lowest cost, best-effort availability, no SLA. Purpose-built for batch processing, model evaluations, and multi-step agentic workflows.

Matching Tiers to Workloads

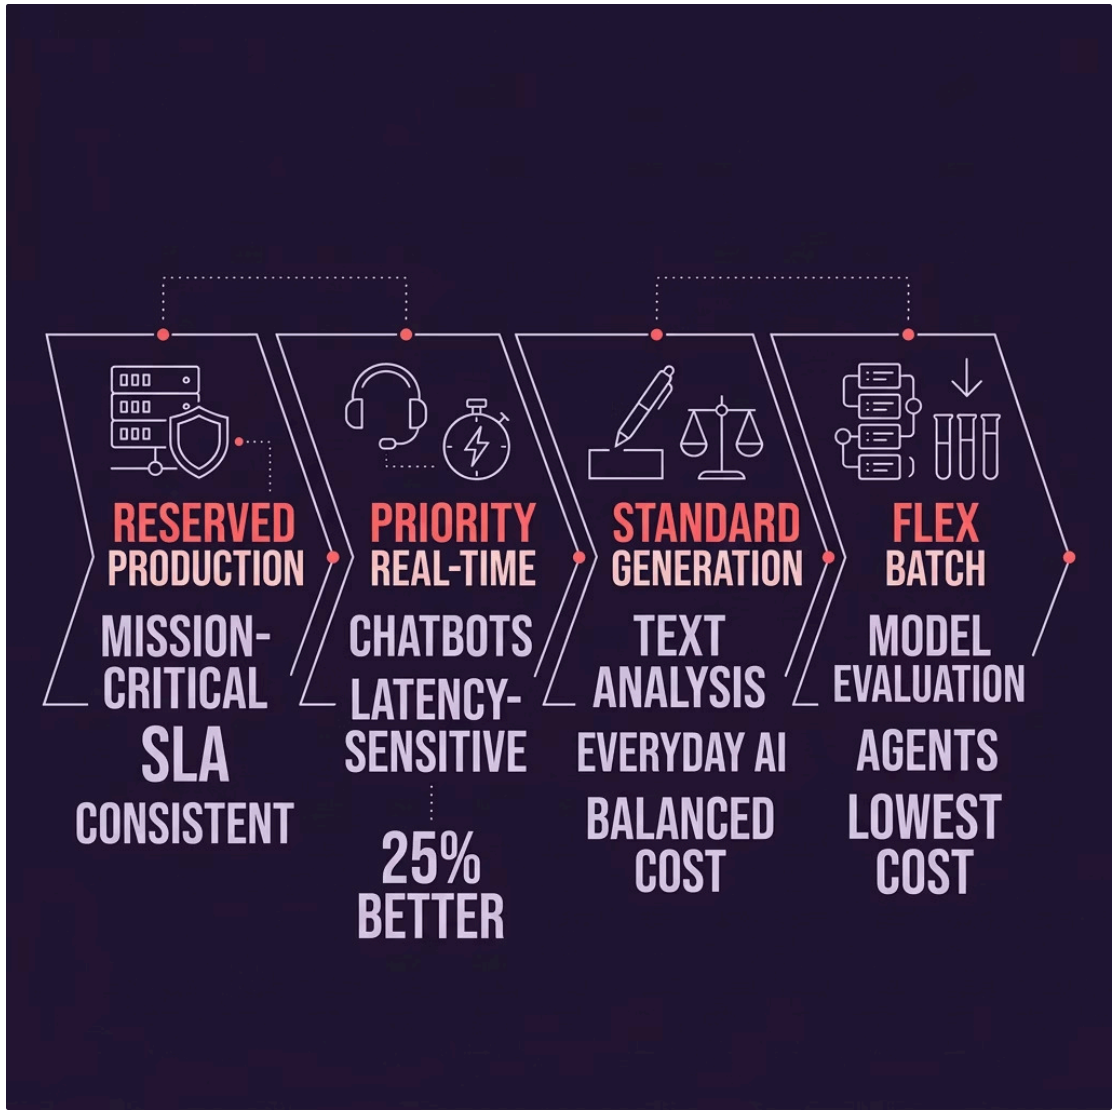

Tier Selection Strategy

Choose your tier based on two primary axes: **latency sensitivity** and **cost tolerance**.

Cost Optimization

Route batch jobs and evaluations to Flex tier — savings can be substantial vs. Standard for non-urgent workloads.

Predictable Capacity

Reserved tier commitments eliminate throttling risk for production workloads with consistent volume requirements.

Hybrid Routing

Combine tiers per use case — Priority for real-time UX, Flex for background pipelines, in the same application.

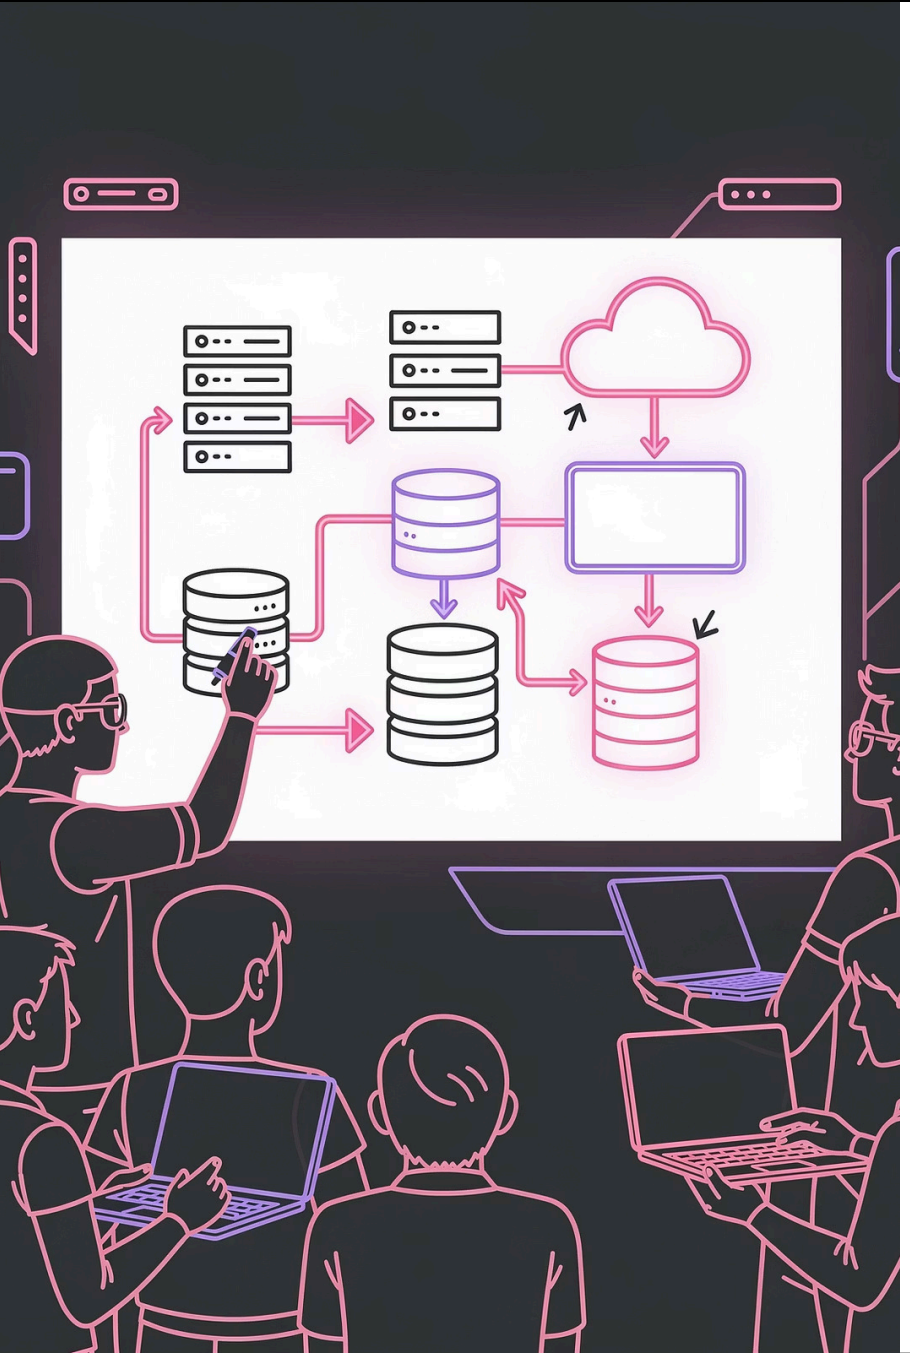

Key Takeaways

1 Decision Framework

Bedrock for managed, production-ready apps; JumpStart when you need custom training, infrastructure control, or specialized ML pipelines.

2 Model Breadth

100+ models from 15+ providers — evaluate across dimensions like context length, modality, licensing, and latency profile before committing.

3 Tier Strategy = Cost Strategy

Map every workload to the right tier: Reserved for SLA guarantees, Priority for UX, Standard as default, Flex for batch and eval jobs. aws.amazon

4 2026 Highlights

OpenAI OSS models, the Bedrock Marketplace, and new service tiers make AWS the most comprehensive managed GenAI platform available today.